

Alexander Polyakov

AI Platform / LLM Systems Engineer - Multi-Agent Workflows · Advanced RAG · MCP · Evaluation

backstabe@gmail.com +995 555 625 979 Telegram: @Aectan LinkedIn: linkedin.com/in/alexanderpolykov

GitHub: github.com/AlexNimoy

Format: Remote · Full-time Experience: 12+ yrs backend · 5+ yrs AI / LLM

Languages: Russian (Native) / English (B2)

ABOUT

AI Platform & LLM Systems Engineer with 12+ years of experience in backend engineering and distributed systems. Specializing in production-ready AI platforms, multi-agent workflows (MCP, supervisor/handoff), and high-performance RAG architectures.

Core focus: • **LLM & Multi-Agent:** autonomous agent workflows, MCP servers, automated evaluation (Ragas, DeepEval, LLM-as-a-Judge, guardrails). • **Advanced RAG:** hybrid search (dense + sparse / BM25), contextual dynamic chunking, reranking, vector DBs (Qdrant, pgvector). • **Engineering & Platform:** async Python (FastAPI, Asyncio), distributed message queues (Kafka, RabbitMQ), System Design, Kubernetes, inference cost optimization.

EXPERIENCE

DFG - Senior AI / LLM Systems Engineer

Nov 2023 – present

USA · dfuture.co · IT, Software Development

- ▶ **AI Tooling & Developer Experience:** Architected and deployed an internal ecosystem of 10+ AI tools, MCP servers, and extensions for Cursor, Claude Code, and OpenCode across 100+ engineers; cut routine development overhead and time-to-market by ~20%.
- ▶ **CI/CD & AI QA Automation:** Configured automated CI/CD pipelines and automated QA workflows using AI for test scenario generation and staging verification, accelerating release cycles and catching regressions early.
- ▶ **Production RAG & Vector Search:** Built and shipped 5+ production RAG pipelines with hybrid search (dense + sparse / BM25) on Qdrant and PostgreSQL (pgvector). Implemented context-aware dynamic chunking and reranking (Cohere Rerank / cross-encoders), reducing latency and increasing Context Precision.
- ▶ **Model Context Protocol (MCP) & Tool Calling:** Designed enterprise system integration layer via MCP, ensuring secure context retrieval and reliable Function Calling / tool use across multi-agent workflows.
- ▶ **LLM Evaluation & Guardrails:** Built automated synthetic evaluation pipeline (Ragas, DeepEval, LLM-as-a-Judge) with Faithfulness, Context Recall, and semantic guardrails, systematically mitigating regressions and hallucinations.
- ▶ **Multi-Agent Orchestration:** Designed agentic workflows (Supervisor, Handoff, stateful graphs) automating complex data analysis and code generation; introduced Spec-Driven Development.
- ▶ **Architectural Leadership:** Authored 20+ ADRs on LLM integration and token spend management; delivered prompt engineering & AI-first workshops for 50+ engineers.

Stack: Python (Asyncio, FastAPI), Golang, MCP, Qdrant, OpenAI / Claude APIs, LangGraph, Ragas, CI/CD, Docker, Kubernetes.

Independent Practice / AI Solutions - Fractional AI & Solutions Architect

Nov 2022 – Nov 2023

B2B & Startup AI Architecture & Engineering

- ▶ **B2B & Startup AI Transformation:** Advised founders and CTOs across 6+ tech companies (FinTech, EdTech, LegalTech) on GenAI system architecture, model selection, and backend integration.
- ▶ **MVP & Production PoC:** Shipped 4+ commercial solutions: intelligent support agents with Function Calling/Tools, semantic search over private knowledge bases, and document processing pipelines.
- ▶ **Cost Optimization (Unit Economics):** Implemented semantic caching (Redis/GPTCache), dynamic model routing (lightweight vs. frontier), and prompt optimization, reducing token costs by 35–50%.
- ▶ **RAG & Vector Databases:** Architected embedding & search pipelines with pgvector and Qdrant; implemented hierarchical chunking and metadata filtering under strict low-latency SLAs.
- ▶ **Technical Audit & Reliability:** Audited API architectures and message queues under peak loads; built rate-limiting, exponential backoff retries, and fallback strategies for provider outages.

Stack: Python (FastAPI, Asyncio), OpenAI API, Anthropic Claude, LangChain, LlamaIndex, pgvector, PostgreSQL, Qdrant, Docker, Redis.

AbsolutTech LLC - Lead Developer / Solution Architect

Nov 2021 – Nov 2022

Moscow, Absolut Insurance Group

- ▶ Launched an insurance anti-fraud product from scratch in ~6 months: requirements, architecture, and production rollout.
- ▶ Architected and deployed a real-time ML anomaly detection pipeline with model inference via FastAPI, RabbitMQ, and Docker on Kubernetes.
- ▶ Evolved system architecture from MVP monolith to 5+ critical microservices in Kubernetes.
- ▶ Led and hired an engineering team of 6: established onboarding, Kanban, and cross-team code reviews.

Stack: Python, FastAPI, AI / ML, Keras, PostgreSQL, Kubernetes, Docker, RabbitMQ, Ruby on Rails 7.

Balance-Platform LLC - Senior Backend / Data Engineer

Feb 2020 – Nov 2021

Moscow, FinTech

- ▶ Engineered a highload distributed credit scoring system integrated with 5+ banks under strict latency SLAs.
- ▶ Built high-throughput feature extraction pipelines and a low-latency statistics aggregation service in Elixir.
- ▶ Optimized complex analytical queries in PostgreSQL and Elasticsearch, resolving queue bottlenecks.

Stack: Python, Elixir, Ruby on Rails, PostgreSQL, RabbitMQ, Elasticsearch, Kubernetes.

Global Search LLC - Backend Engineer

Sep 2018 – Jan 2020

Saint Petersburg

- ▶ Developed core modules of an insurance aggregator: resilient billing with automated payouts and fraud prevention components.
- ▶ Optimized PostgreSQL transactional performance and configured automated CI/CD pipelines, halving release cycles.

Stack: Ruby on Rails, PostgreSQL, Redis, Docker, CI/CD.

Go-promo / Neuron digital - Backend & Data Engineer

May 2016 – Feb 2018

Moscow · neuron.digital

- ▶ Built high-performance backend services and ETL pipelines for high-traffic media platforms serving millions of users.
- ▶ Implemented analytical warehouse and search using ClickHouse and Elasticsearch; built async event bus via RabbitMQ.

Stack: Python, Golang, PostgreSQL, ClickHouse, Elasticsearch, RabbitMQ, Docker.

PolexGroup - Software Engineer

Jan 2014 – May 2016

Moscow · polexgroup.ru

- ▶ Developed ERP/CRM modules and engineering calculation algorithm services for internal operations.
- ▶ Designed relational schemas and optimized complex SQL queries in PostgreSQL.

Stack: Python, PostgreSQL, Linux, Git.

SKILLS

AI / LLM: Python · LangGraph · MCP · RAG · hybrid search · LlamaIndex · LangChain · LLM Evaluation · Ragas · DeepEval · LLM-as-a-Judge · guardrails · Qdrant · vector search · embeddings · prompt engineering · function calling / tool use · multi-agent orchestration (supervisor / handoff) · Langfuse · LiteLLM · vLLM · OpenAI / Claude APIs

Backend & Architecture: Asyncio · FastAPI · Pydantic · Golang · System Design · SQLAlchemy · SQL · Ruby on Rails · Elixir

Data & Messaging: PostgreSQL · pgvector · Qdrant · Apache Kafka · RabbitMQ · Redis · Elasticsearch · ClickHouse

DevOps & ML: Kubernetes · Docker · CI/CD · Linux · Git · PyTorch · pandas · NumPy · team leadership · solution architecture

EDUCATION

Moscow Financial and Industrial University (Synergy), 2025 - Higher Education: Information Systems & Software Engineering, Information Systems Software Development and Information Security.

Alexander Polyakov · backstabe@gmail.com · +995 555 625 979 · Telegram @Aectan · linkedin.com/in/alexanderpolykov · github.com/AlexNimoy